A POINT OF VIEW · AGENTIC AI IN BANKING

Agents that change the *work*, not just the worker.

A primer on agentic AI in banking – and why the next decade of value comes from redesigning processes, not buying better assistants.

PRACTICE	EDITION	AUDIENCE	DISTRIBUTION
Banking	*No. 01 · May 2026*	*CXO buyers · IC pursuit teams*	*Sales enablement – Internal*

CHAPTER – · TABLE OF CONTENTS

Six chapters and an appendix. About 25 minutes if read in sequence.

Three things to take from this deck

CHAPTER I · SETUP

The argument, in three lines.

01

THE TRAP

The median bank runs 40+ generative AI pilots; realized value concentrates in only two or three. The pilot factory has stopped paying off.

02

THE REFRAME

Pilots plateau because agents inherit the friction of the process they joined. The prize is redesigning the process around what an agent can actually do.

03

THE PATH

Five no-regret moves now, one redesigned process within twelve months, and an operating-model shift across the next twenty-four.

CHAPTER I · WHY NOW

A phase change, not a headline.

i.

Capability

Frontier models can plan, use tools, and verify their own work. Per-token cost has fallen ~85% since 2023, putting capable models in reach of bulk operational workloads.

ANTHROPIC, OPENAI, GOOGLE · 2024-2026

ii.

Tooling

Orchestration, evaluation harnesses, observability, and guardrails have moved from research-interesting to production-engineering. Banks can buy or self-host an agent stack like they buy analytics.

LANGGRAPH · OPENAI ASSISTANTS · ANTHROPIC

iii.

Maturity

JPMorgan reports $17B+ in annual tech spend with AI as a named line item. BofA's Erica has logged 2B+ client interactions. The infrastructure is in place; the design discipline is not.

JPM 10-K 2024 · BCG BUILD FOR THE FUTURE 2024

CHAPTER I · THE TRAP

The pilot factory has *stopped* paying off.

Across the largest banks, a long-tail distribution: a small number of use cases deliver concentrated value while most pilots remain stuck in proof-of-concept.

FIG. 01

Pilot value concentrates in a long tail.

Per-bank distribution of realized value across active generative-AI initiatives. The top two carry the majority.

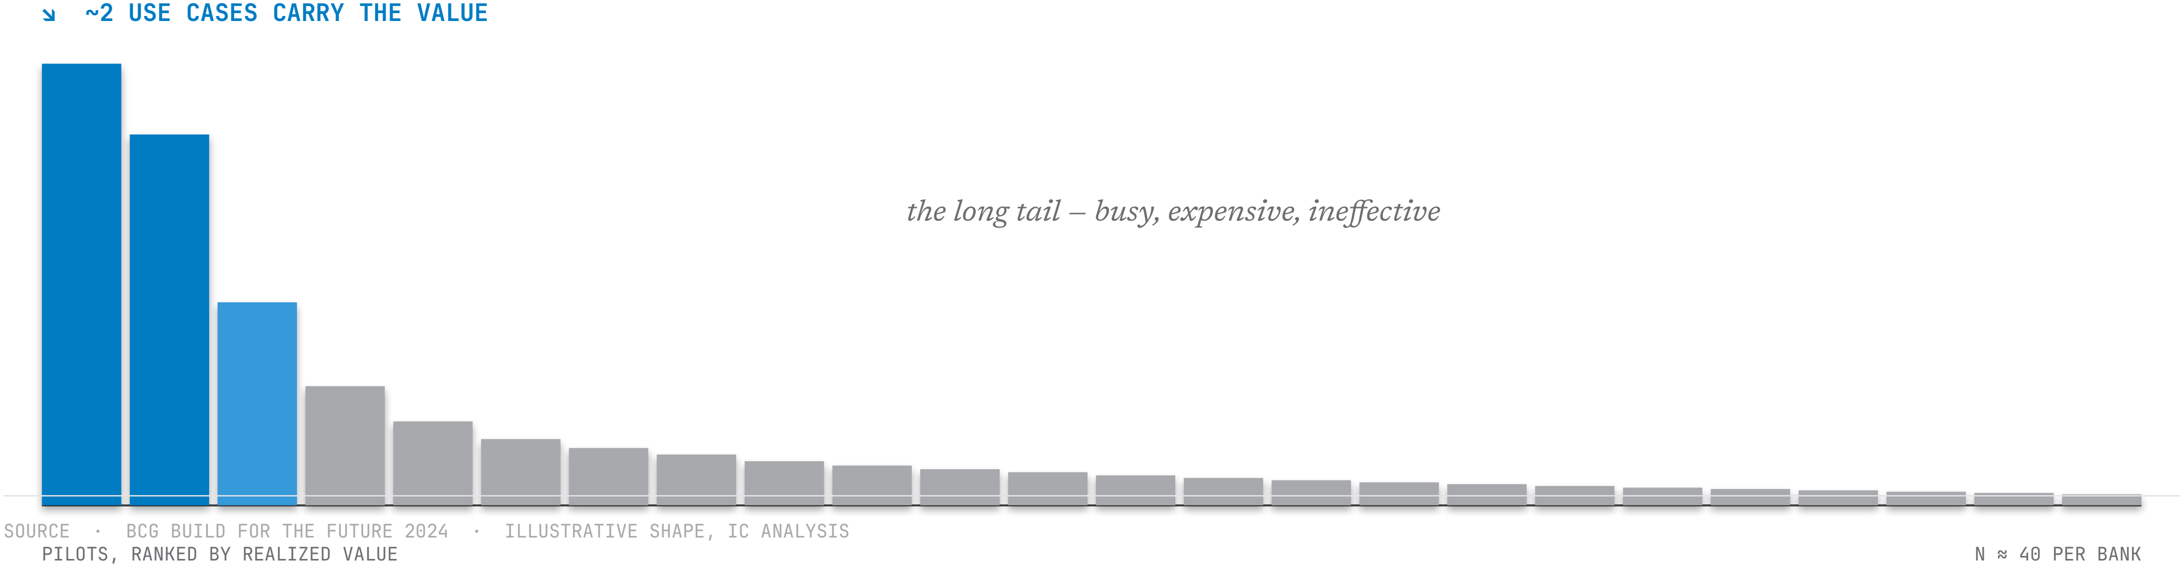

SOURCE · BCG BUILD FOR THE FUTURE 2024 · ILLUSTRATIVE SHAPE, IC ANALYSIS

CHAPTER II · REFRAME · OUR POINT OF VIEW

Banks have been adding agents to processes. The next decade is about *redesigning processes around agents.*

Pilots plateau because the agent inherits the friction of the workflow it joined – the handoffs, the data silos, the compliance choke-points, the human approvals built for a slower world. The teams who break out aren't the ones with better models. They are the ones who redrew the process around what an agent can actually do, and rebuilt the operating model – control, risk, talent – to match.

CHAPTER II · PILLARS

Six pillars define the move from agent-augmented to agent-native.

i.

Control plane

Orchestration, routing, retries, escalation. The agent's wiring is a first-class system.

ii.

Agent roles

Each agent has one clear job – research, decide, act, or review. Roles compose; sprawl is avoided.

iii.

Human-in-the-loop

Humans handle exceptions and decisions of consequence. They sit at the seam, not at the start.

iv.

Observability

Every step is logged, replayable, and scored. What you cannot measure, you cannot improve.

v.

Governance

Model risk, audit trail, regulator-readable explanations are designed in, not bolted on.

vi.

Talent

New operating roles: agent designer, agent-ops, agent product manager. The org chart shifts.

The bank, redrawn: where the unit of work changes

CHAPTER II · HOT-SPOTS

A finite set of processes are candidates for agent-native redesign. The blue marks show where the work itself changes shape.

FRONT OFFICE

Customer, advisor, channel-facing

- Relationship banker copilot
- Wealth advisor co-pilot
- Digital channel agentic flows
- Contact-centre resolution agent

MIDDLE OFFICE

Risk, ops, control

- KYC & onboarding orchestration
- Credit underwriting & memo
- Dispute handling & servicing
- Fraud investigation co-pilot

BACK OFFICE

Finance, ops, regulatory

- Reconciliation & break-resolution
- Regulatory reporting drafting
- Treasury & liquidity ops
- Audit & compliance evidence

CHAPTER III · THREE PATTERNS

14 systems, 11 handoffs, 9 days — and the analyst spent most of it moving documents.

TODAY

The legacy process

- ~14 systems, 11 handoffs
- 9-day median TAT for commercial
- 8–12% drop-off in customer journey
- Analyst-led, document-by-document review

SEAM

REDESIGNED

The agent-native version

+ Single orchestrated agent flow
+ Documents read, profiles drafted, risk tiered
+ Humans approve only mid/high-risk cases
+ Audit trail generated continuously

REDESIGN PRINCIPLES

The customer enters the system once. The system does the joining.

Documents go to an agent first, an analyst second.

Control sits in the orchestration layer, not the screens.

The memo is generated, not authored. The reviewer sees the working.

TODAY

The legacy process

- Analyst pulls statements, spreads, market data
- Manual memo drafting over weeks
- Limited scenario coverage (2–3 cases)
- Reviewer sees the answer, rarely the working

SEAM

REDESIGNED

The agent-native version

+ Research agent ingests filings + market data
+ Decision-support agent runs spread and scenarios
+ Human reviewer at the seam, not the start
+ Memo generated from the agent's working

REDESIGN PRINCIPLES

The memo is generated, not authored.

Scenarios are cheap; run more of them.

Reviewers see the working, not just the answer.

CHAPTER III · THREE PATTERNS

The agent owns the case. The human owns the exception.

TODAY

The legacy process

- 25-day median resolution cycle
- Customer re-tells story to 2–3 teams
- Ledger entries lag the case state
- Communications are ticket-driven

SEAM

REDESIGNED

The agent-native version

+ One agent owns the case end-to-end
+ Ledger is the source of truth; screens are the lens
+ Humans handle exceptions only
+ Communications generated proactively

REDESIGN PRINCIPLES

The agent owns the case; the human owns the exception.

The ledger is the truth; the screen is the lens.

Communications are continuous, not on request.

CHAPTER III · SYNTHESIS

Five recurring principles, regardless of which process you start with.

I	Events, not handoffs.	The process becomes an event stream. State is shared, not passed between teams.
II	Verify, don't trust.	Every agent output is scored against an eval before it advances. No eval, no deployment.
III	Humans at the seam.	Humans handle exceptions and decisions of consequence – not data movement or formatting.
IV	Ledger first, screen second.	The system of record drives the experience. The screen is a view, not the truth.
V	One agent, many tools.	Agents specialize by role, not by screen. Capability composes; sprawl is avoided.

CHAPTER IV · MATURITY

Maturity is not a model question. It is an operating-model question.

STAGE I

· TODAY ·

Pilots

One-off use cases. Manual scaling. Value rarely attributed back to a P&L.

STAGE II

· TODAY ·

Patterned use cases

Templates emerge. Platform investments begin. Evals are introduced.

STAGE III

· THE PRIZE ·

Agent-native processes

Workflows redesigned end-to-end. KPIs rewritten. Operating model adjusted.

STAGE IV

· THE FUTURE ·

Agent-native bank

Roles, P&L, and controls are built around agents. The org chart reflects it.

CHAPTER IV · TODAY

Most large banks remain between stages one and two. The gap to stage three is where the next decade of advantage is decided.

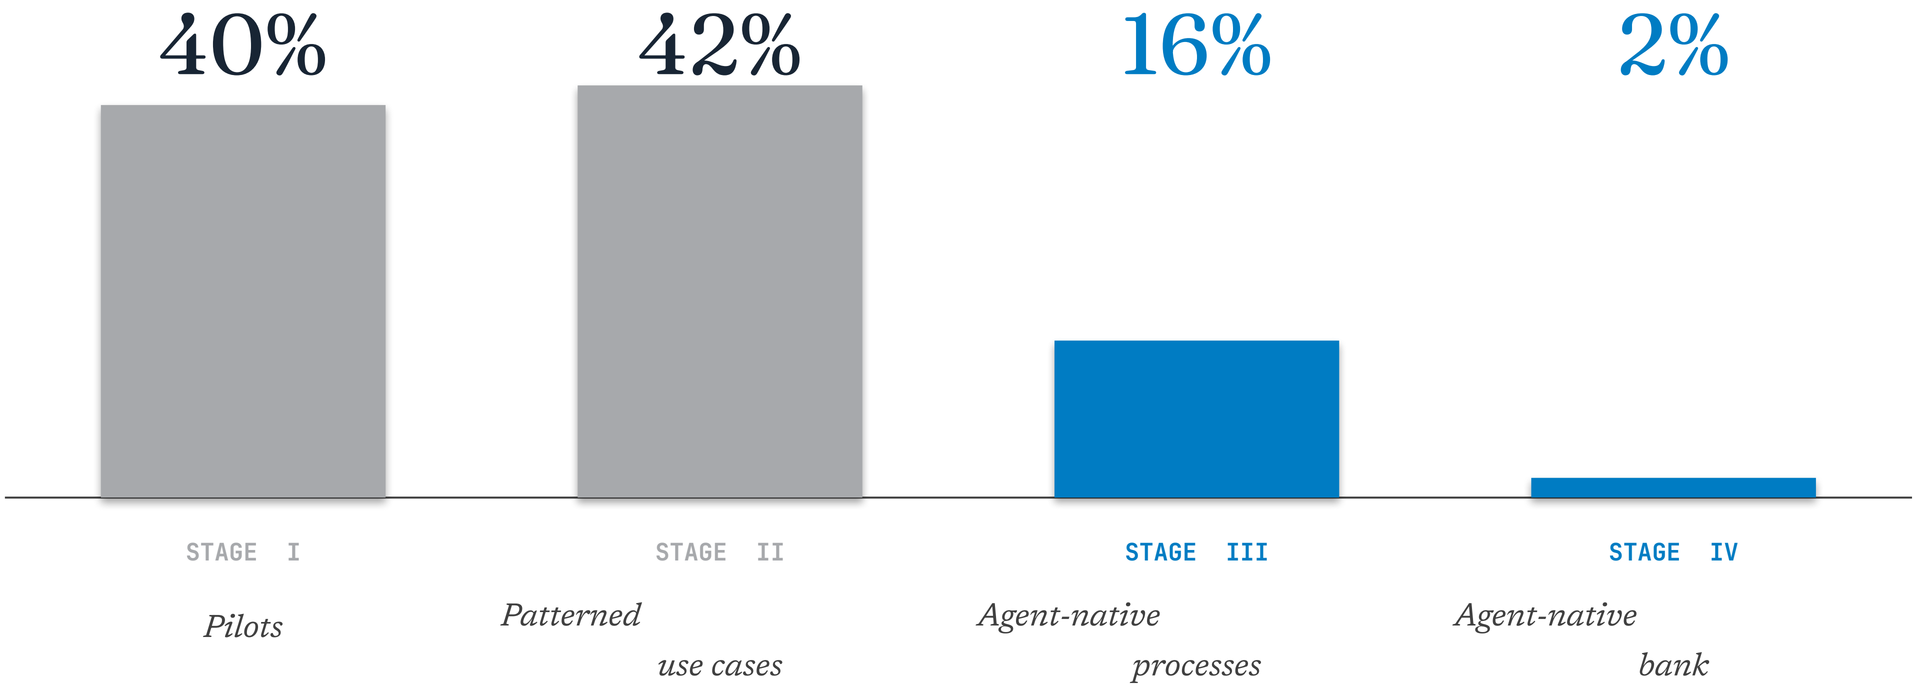

The gap between stages II and III is where the next decade is decided.

SOURCE · IC ANALYSIS OF PUBLIC DISCLOSURES, MCKINSEY STATE OF AI 2024, BCG 2024 · ILLUSTRATIVE

The path to value: what to do across 24 months

CHAPTER IV · THE PATH

Three horizons. Leaders cross from pilots to a live agent-native process inside this window.

0 – 3 MONTHS

Build the case

- Pick 1–2 candidate processes for redesign
- Quantify the value at stake; secure executive sponsor
- Stand up evals + governance
- Recruit / re-skill the core agent team

3 – 12 MONTHS

Ship one process

- Redesign and deliver one process end-to-end
- Codify reusable patterns
- Capture and attribute savings
- Begin shifting KPIs around the new process

12 – 24 MONTHS

Scale the model

- Roll patterns across the function
- Rewrite productivity and risk KPIs
- Restructure ops around agent-native flows
- Stand up agent-ops as a permanent capability

Five no-regret moves to start now

CHAPTER IV · MOVES

Infrastructural, not speculative. They pay back in every plausible future state of the technology.

i **Establish an agent control plane.**

Pick one orchestration stack; standardize. Multiple stacks fragment evals, controls, and talent.

ii **Build the eval discipline.**

Every agentic capability ships with a measurable scorecard. No eval, no promotion.

iii **Pick a redesign, not a copilot.**

Choose a process worth redrawing – not just augmenting. Copilots have known ceilings; redesigns don't.

iv **Get the governance ready early.**

Model risk, audit trail, regulator-readable explanations from day one. First-movers shape the precedent.

v **Reshape the talent contract.**

Hire and train agent designers, agent-ops, and agent product managers. The role gap closes in 6–12 months.

PLACEHOLDER · REPLACE WITH IC OFFICIAL COPY

Four capabilities, across the agentic lifecycle.

STRATEGY

Shape the bet

- Agentic AI strategy & value-pool mapping
- Operating-model design
- Investment case & ROI framing

DESIGN

Redraw the process

- Process redesign for agent-native flows
- Agent role & control-plane design
- Governance & risk framework

BUILD

Make it real

- Agent engineering & orchestration
- Integration with core banking
- Evals, observability, guardrails

RUN

Keep it improving

- Agent-ops capability stand-up
- Model risk & continuous monitoring
- Capability rollout & change

CHAPTER V · STACK

PLACEHOLDER · CONFIRM IC POSITIONING

The agent stack — and the layers where IC delivers.

IC PLAYS HERE	*Change management & adoption*	Operating-model change, role redesign, learning
IC PLAYS HERE	*Governance, model risk & audit*	Risk framework, audit trail, regulator readiness
IC PLAYS HERE	*Agent operations*	Run, monitor, improve, retire – the lifecycle
IC PLAYS HERE	*Integration with core banking*	Connect agents to systems of record
IC PLAYS HERE	*Agent design & orchestration*	Roles, control plane, evals, guardrails
ADVISORY ONLY	*Tooling & infrastructure*	Frameworks, vector stores, observability
ADVISORY ONLY	*Foundation models*	Models from providers (selection advisory only)

Qualifying questions for client conversations

CHAPTER V · TOOL 01

Ten questions, four dimensions. The answers reveal where the client sits.

AMBITION

Q. *What process would you redesign if you could start over today?*

Q. *If agents become 50% cheaper next year, what changes in your operating model?*

CURRENT STATE

Q. *How many of your pilots have measurable value attribution today?*

Q. *Which use cases are in production with daily users?*

Q. *Which use cases were quietly retired in the last 12 months?*

RISK APPETITE

Q. *What level of agent autonomy is your regulator comfortable with today?*

Q. *Which decisions must always have a human reviewer – and why?*

OPERATING-MODEL READINESS

Q. *Who owns 'agent ops' in your organization?*

Q. *How will you measure productivity in an agent-native team?*

Q. *Where will the savings show up – and who is accountable for them?*

Common objections — and how to answer them

CHAPTER V · TOOL 02

Five things you will hear in the first 30 minutes — and how to answer each.

"Our regulator will never allow autonomous agents."	↳	Start with agent-assisted, not autonomous. Design the audit trail before the agent. Regulators react to surprises, not to readable explanations.
"We need to wait for the tech to stabilize."	↳	The agent layer changes; the redesign value is durable. Build patterns, not POCs.
"We've tried — pilots don't scale."	↳	Pilots fail because the process is unchanged. Redesign the process, then scale. The plateau is a design problem, not a model problem.
"Our data isn't ready."	↳	Agents tolerate messier data than traditional ML. Start with read-mostly use cases. Data investment runs in parallel.
"We don't have the talent."	↳	Talent reshapes faster than tech. The role gap is solvable in 6–12 months with a mix of hires and re-skilling.

Illustrative case studies

CHAPTER V · CASES

PLACEHOLDER · REPLACE WITH REAL ENGAGEMENTS

TOP-10 US BANK

Agent-native KYC

SITUATION

14-system onboarding, 9-day TAT, 8% drop-off

REDESIGN

Orchestrated agent flow with risk-tiered checkpoint

OUTCOME

60%+ TAT reduction, 40% analyst capacity returned

EUROPEAN UNIVERSAL BANK

Credit underwriting copilot

SITUATION

Multi-week memo cycle, 2–3 scenarios per deal

REDESIGN

Research + decision-support agents, reviewer at the seam

OUTCOME

35% memo turnaround reduction, 2× scenario coverage

REGIONAL COMMERCIAL BANK

Dispute handling redesign

SITUATION

25-day cycle, customer re-tells story to multiple teams

REDESIGN

Single agent owns case end-to-end, ledger-first

OUTCOME

25-day → 7-day median resolution, +18 NPS lift

ALL THREE CASES ILLUSTRATIVE · REPLACE WITH REDACTED REAL IC ENGAGEMENTS BEFORE EXTERNAL USE

Stop adding agents. Start *redesigning* around them.

01 The pilot plateau is a design problem, not a model problem.

02 Value lives in process redesign and operating-model change – not in better assistants.

03 The next 24 months will separate banks that 'use AI' from banks that are agent-native.

PLACEHOLDER · CONFIRM IC CONTACT DETAILS

SUGGESTED NEXT STEPS

Operating-model workshop

A 2-day session with the executive sponsor to identify the first redesign candidate.

Maturity assessment

A 3-week diagnostic mapping current state against the four-stage maturity model.

Value pool diagnostic

A 4–6 week analysis of where agent-native redesign moves the needle most.

TALK TO US

[Name]

[Title – Banking practice, Infosys Consulting]

[email@infosysconsulting.com]

[+1 ___ ___ ____]

Or ask your IC account team to set up a working session.

— CHAPTER END

Appendix.
Sources & *references.*

CHAPTER — · APPENDIX

Supporting research and citations across the deck.

MARKET SIZING & MATURITY

- McKinsey & Co. The economic potential of generative AI, June 2023.
- McKinsey & Co. The state of AI in early 2024 – global survey.
- Boston Consulting Group. Build for the Future 2024.
- Gartner. GenAI implementation survey 2024 – pilot abandonment.

BANKING-SPECIFIC EVIDENCE

- JPMorgan Chase. 2024 Annual Report and 10-K filings.
- Bank of America. Investor day disclosures (Erica interactions).
- Citigroup. Public statements on agentic AI initiatives, 2024-2025.
- Goldman Sachs. Public disclosures on internal AI deployments.

PROCESS BENCHMARKS

- Celent. KYC and onboarding benchmark research.
- Forrester. Onboarding and dispute-handling benchmarks.
- Aite-Novarica. Commercial onboarding TAT industry data.

GOVERNANCE & REGULATION

- Federal Reserve. SR 11-7: Guidance on Model Risk Management.
- OCC. Heightened Standards for Large Banks.
- European Union. Artificial Intelligence Act, in force 2024.
- UK FCA. AI principles and discussion papers.